

In this part, you will implement another web server – HTTPAsk. This is a web server that uses the client you did in Task 1 and 2. When you send an HTTP request to HTTPAsk, you provide the parameters for TCPClient as parameters in the HTTP request.

When HTTPAsk receives the HTTP request, it will call the method `TCPClient.askServer()`, and return the output as an HTTP response. This may seem confusing, but it is really very simple: instead of running TCPAsk from the command line, you will build a web server that runs TCPAsk for you, and presents the result as a web page (in an HTTP response).

What you will learn here is to:

- Read and analyse an HTTP GET request, and extract a query string from it.
- Compose and return a valid HTTP response.

Task

Your job is to implement a class called HTTPAsk. It's "main" method implements the server. It takes one argument: the port number. So if you want your server to run on port 8888, you would start it in the following way:

```
$ java HTTPAsk 8888
```

You will be provided with a file with a template for the HTTPAsk class. There is not much content in it though, but your task is to fill in the missing Java code in *this file*.

Since HTTPAsk is a TCP server, you should implement it in the same way as other servers: it should run in an infinite loop, and when one client has been served, the server should be prepared to take care of the next. Your server does not need to serve more than one client at a time, though. It does not need to be multi-threaded, in other words (that's for later).

Instructions and Tips

- Use the `ServerSocket` class to create the socket for the server.
- Keep in mind that you are required to use byte streams I/O with explicit encoding/decoding. No wrapper classes!
- It is tricky to get everything right at once. Follow a sound software engineering approach and develop your code step by step. After each step, design and run one or more tests to verify that your code works.

Start with the HTTP Server

A suitable first step can be to create an HTTP server that generates a static HTTP response. In other words, when a client connects to the server, ignore the data that the client is sending. Instead, just return an HTTP response that always has the same data (like "hello world"). Then use a browser to connect to the server, and check that the response is shown correctly in the browser. In this way, you can verify that you can generate a correct HTTP response that the browser recognises.

Some things to consider when generating HTTP messages.

- Line breaks are significant in HTTP: Each field in the header (header line) is terminated with a line break. The header is separated from the body by an empty line (= two consecutive line breaks.)
- A line break in HTTP consists of two characters: carriage return followed by line feed. This is probably not what your operating system uses to mark end of line (it probably uses only a line feed). In Java, the syntax for carriage return followed by line feed in a string is "\r\n". So make sure to write this sequence to mark the end of each line.
- A web browser understands and can present a plethora of different data formats (the browser often tells the server what formats it recognises through an "Accept" header line in the HTTP request). So what format should you use? Luckily, the default format for HTTP is plain text, "text/plain". As you can tell by the name, this format is just plain text, without any formatting. So unless your server specifies something else in its HTTP response, the web browser will assume your response is in plain text, and present it as such in the browser window. (If you want, you can format the response more nicely with HTML code, but that is not a requirement. But if you do that, you should specify that the content format is HTML – with the header line "Content-Type: text/html".)
- An HTTP response should also specify the character encoding that is used. The default is UTF-8, which is a superset of ASCII, so here you can just use the default as well. In other words, you don't need to specify this either.

Parsing the HTTP Request

What you need to work on in this task is mainly to process the HTTP request. The HTTPAsk server recognises one "resource" in the URL, called "ask". So the general format of an HTTP request for HTTPAsk is as follows:

```
http://hostname.domain/ask?<parameters>
```

Here, `<parameters>` represents the arguments for the `TCPClient` instance.

You should extract the hostname, port, and string parameters from the request, as well as the optional arguments for `askServer` (ie., shutdown flag, timeout, and limit). Your HTTP server only needs to recognise GET requests. In a GET request, parameters are sent in the URI component (the second component) of the first line of the request. The URI component comes from the URL that was used to navigate to the server.

So, for instance, this is the URL you would type into the navigation field of your browser to tell the server to access the "daytime" service at "time.nist.gov":

```
http://hostname.domain/ask?hostname=time.nist.gov&limit=1200&port=13
```

The browser will then open a TCP connection to "time.nist.gov" at port 13, and send the following GET request over the connection:

```
GET /ask?hostname=time.nist.gov&limit=1200&port=13 HTTP/1.1
```

```
Host: hostname.domain
```

Note that there is an empty line (two consecutive line breaks) that marks the end of the GET request header. last The "/ask" part of the URI tells the HTTPAsk server to set up a TCP connection to a remote server (specified by the "hostname" parameter) and a port ("port" parameter), with a data limit of 1200 bytes, and return the result. Think of this as a command to the server!

You need to write the code that validates the request, "picks apart" the URI component of the request and extracts the parameters and their values. This involves a certain amount of string processing. There are (at least) two ways you could do this. One way is to use the [URL class](#), which has many different methods for processing URL strings. Another way is to take care of the string processing yourself – you will probably find that the [String "split" method](#) is very handy for this. It is up to you!

The query parameters your HTTPAsk server should recognize are shown in the following table:

| Parameter | Meaning | Example |
|-----------|---|-----------------------|
| hostname | Domain name for host | hostname=whois.iis.se |
| port | Port number | port=43 |
| string | Data to send to server Optional Note: before passing the string to <code>askServer</code> (as a byte array), make sure to terminate the string with a line break "\n". | string=kth.se |
| shutdown | Shutdown flag Optional, default is false | shutdown=true |
| limit | Data limit in bytes Optional, default is no limit | limit=512 |
| timeout | Maximum wait time in milliseconds Optional, default is no timeout | timeout=2000 |